

# Bijjective Mapping Invariant

The bijective mapping invariant of the Theseus OS states:

*Each page in the virtual address space of the system can only be mapped to one frame of the physical address space, and vice versa.*

**Theseus Memory Subsystem Background:** In Theseus, we create a **Frames** and **Pages** type, instances of which are the unique representation of a region of physical or virtual memory, respectively. With `typestate` programming, we create four possible states for instances of the **Frames** and **Pages** types: **Free**, **Allocated**, **Mapped**, and **Unmapped**.

The memory subsystem of Theseus only adds a Page Table Entry (PTE) when creating a **Pages<Mapped>** instance, and only removes it when dropping the same instance. Thus, proving the bijective mapping invariant necessitates proving the correct construction and destruction of a **Pages<Mapped>** instance.

**Proof Structure:** We list the Lemmas required to prove the invariant, and next to each lemma we list the techniques used to prove it: formal verification (**F**), the type system (**T**), or manual checks (**M**). We use a prose proof (**P**) to tie multiple techniques together.

## Prose Proof:

- ⟨1⟩1. **Frames** and **Pages** instances are unique. [**F,T,M,P**]
- ⟨2⟩1. **Chunk** instances created by a single allocator are unique. [**F,T,M,P**]
- ⟨3⟩1. **Chunk** instances cannot be cloned or copied as **Chunk** is a linear type and doesn't implement the **Clone** trait. [**T**]
- ⟨3⟩2. The inner range value of a **Chunk** instance is not changed in unverified functions. [**T,M,P**]
- ⟨4⟩1. Visibility modifiers, and preventing implementation of the **DerefMut** trait, limit access of the range value to **Chunk** methods only. [**T**]
- ⟨4⟩2. Manual checks of unverified methods show they only take immutable access to the **Chunk**. [**M**]
- ⟨4⟩3. *Necessary and Sufficient:* These conditions ensure that there is no way to mutably access the inner range value except through verified functions
- ⟨3⟩3. All methods that create or mutate a **Chunk** instance are proven to forbid overlapping ranges. [**F,T,M,P**]

- (4)1. The `new()` method is formally verified to only create a **Chunk** instance if its range value does not overlap with all previously instantiated **Chunks**, information about which is stored in a verified data structure. When a new **Chunk** is instantiated, its information is also added to the verified data structure. [F]
- (4)2. The `new_from_pte()` method creates a **Chunk** for a range of frames that have just been unmapped. It is only used by the frame allocator. [F,T,M,P]
  - (5)1. This function is formally verified to create a **Chunk** from the range value that it is passed. [F]
  - (5)2. The range value that is passed is the inner value stored in **UnmapResult**. [M]
  - (5)3. **UnmapResult** is a linear type that's only created as the output of a trusted PTE HAL function. [T,M,P]
  - (5)4. *Necessary and Sufficient:* **UnmapResult** is a linear type that acts as a proof of work, that the range it stores represents the frames that have been unmapped from the page table and a representation doesn't exist in the system. `new_from_pte()` only takes this range as an argument to recreate a **Chunk**.
- (4)3. The `merge()` method merges the ranges of two contiguous **Chunks** to create one **Chunk** with the same range. [F,T,P]
  - (5)1. `merge()` is formally verified to combine the range of a **Chunk** instance with the range of a contiguous **Chunk**. If the arguments passed are not contiguous, both **Chunk** instances are unchanged. [F]
  - (5)2. `merge()` consumes and drops the **Chunk** that is merged, so it no longer exists at the end of the function. [T]
  - (5)3. *Necessary and Sufficient:* `merge()` combines two **Chunk** instances into one, and prevents any overlapping **Chunk** from existing at the end of the function.
- (4)4. The `split_range()` method takes one **Chunk** and splits it into 1-3 **Chunks**. [F,T,P]
  - (5)1. `split_range()` is formally verified to split a **Chunk** instance into 1-3 **Chunks** depending on the range passed as an argument. If the range to be extracted does not lie within the range, the **Chunk** instance is unchanged. [F]
  - (5)2. `split_range()` consumes and drops the **Chunk** that is split, so it no longer exists at the end of the function. [T]
  - (5)3. *Necessary and Sufficient:* `split_range()` prevents any overlapping **Chunks** from existing at the end of the function.
- (4)5. The `split_at()` method takes one **Chunk** and splits it into two. [F,T,P]
  - (5)1. `split_at()` is formally verified to split the range of a **Chunk** instance at a given offset and create two new **Chunk** instances. If the offset does not lie within the range, the **Chunk** instance is unchanged. [F]

- (5)2. `split_at()` consumes and drops the `Chunk` that is split, so it no longer exists at the end of the function. [T]
  - (5)3. *Necessary and Sufficient*: `split_at()` prevents any overlapping `Chunk` from existing at the end of the function.
  - (4)6. *Necessary and Sufficient*: These are the only means by which a `Chunk` can be instantiated or mutated. They ensure that it is unique at the time of creation, and it is only mutated if we've proven that no other `Chunk` instance with an overlapping range exists (by consuming the `Chunk` instance that has the overlapping range)
- (3)4. *Necessary and Sufficient*: These conditions prove that a `Chunk` is unique at the time of instantiation. Then for its entire lifetime, it can never be duplicated or mutated in a way that makes it lose its uniqueness.
- (2)2. `Frames` and `Pages` are composed of the `Chunk` type. [T]
- (2)3. `Frames` and `Pages` cannot be copied or cloned. The Rust type system ensures that a type that owns a linear type is linear so as `Chunk` does not implement the `Clone` trait, neither do `Frames` or `Pages`. [T]
- (2)4. The range value of a `Chunk` instance stored in a `Frames` or `Pages` cannot be directly mutably accessed. Type system visibility modifiers keep the `Chunk` type opaque. [T]
- (2)5. *Necessary and Sufficient*: These conditions ensure that the linearly-typed fields of the `Frames/ Pages` type are unique at the time of instantiation which means the `Frames/ Pages` instance is unique. After instantiation, a `Frames/ Pages` cannot be duplicated or mutated in such a way that could make it lose its uniqueness. The field that is unique remains unique for the lifetime of the instance.
- (1)2. A `Page<Mapped>` instance can only be created by taking ownership of both a `Pages<Allocated>` instance and a `Frames<Allocated>` instance. The mapping interface is defined as `map(frames: Frames<Allocated>, pages: Pages<Allocated>, ...) -> Pages<Mapped>`. [T]
- (1)3. PTEs can only be manipulated through the memory subsystem's `map` and `unmap` functions. [T,M,P]
  - (2)1. Visibility modifiers restrict access to functions that manipulate the page table to within the memory crate. [T]
  - (2)2. Manual checks within the memory crate ensure that PTEs are only added in the `map` and `unmap` functions. [M]
  - (2)3. *Necessary and Sufficient*: the type system makes manual checks easier by limiting where they have to occur. Manual checks make sure that functions are only used in the correct places.
- (1)4. Creating a `Pages<Mapped>` instance adds PTEs only for the given pages and frames, and dropping it removes them. Manually check 10 lines of sequential code in the `map/ unmap` functions. [M]
- (1)5. *Necessary and Sufficient*: Lemma 1 ensures that, at the time of creation of a `Pages<Mapped>`, the representation of the pages and frames that are about to be mapped are unique. Lemma 2 ensures that the pages/frames cannot be used for any other mapping while this mapping exists. Lemma

3 ensures that a mapping cannot be created in a way that circumvents the guarantees provided by **Pages<Mapped>**. Lemma 4 ensures that software state accurately reflects hardware state, such that the compiler can enforce invariants about page table contents.